PRODUCT MANAGEMENT CASE STUDY

Reducing User Dissatisfaction at Dubizzle

A Review-Driven Product Analysis | UAE Classifieds Platform

Analysis Period: 2023 – 2026 | **Data Source:** Google Play Store User Reviews (884 reviews)

884	70.1%	3	74%
Reviews Analyzed	Negative (1–2 Star)	Critical Problem Clusters	Overall Response Rate

01 | Context & Background

Dubizzle is the UAE's dominant classifieds platform, operating across real estate, motors, jobs, and general goods. Originally launched in 2005, it has a large and established user base but faces growing competition from Facebook Marketplace, Opensooq, and vertical-specific platforms such as YallaMotor and Bayut.

This case study simulates a product analysis exercise conducted as part of a self-initiated portfolio project. The objective was to practice core PM skills—user research synthesis, problem prioritization, and solution framing—using real, publicly available data from Google Play Store reviews.

No internal product data, engineering resources, or team access were involved. All findings and recommendations are grounded solely in the review dataset.

02 | Problem Statement

With 70.1% of the analyzed reviews rated 1 or 2 stars, Dubizzle is experiencing meaningful dissatisfaction among its most vocal users. While a vocal minority of reviewers skews any public dataset negative, the sheer volume and consistency of specific complaints across hundreds of reviews signals real friction that is likely affecting retention, posting behavior, and advertiser trust.

The core question driving this analysis was:

> *"Which recurring user pain points, if addressed, would most meaningfully improve the experience for the largest number of Dubizzle users—without requiring an overhaul of the entire product?"*

03 | Research Methodology

Approach

The dataset consisted of 884 Google Play Store reviews across four rating buckets (1-star through 4-star). The analysis followed a structured three-step process:

- **Step 1.** Data Extraction & Structuring:
  - Parsed the raw Word document into a structured Excel dataset with columns for reviewer name, review text, star rating, date, and Dubizzle's reply (if any).
  - Confirmed reply/no-reply status for each review to analyze response coverage.
- **Step 2.** Thematic Coding:
  - Applied keyword-based pattern matching across the review text to tag each review with one or more pain point categories.
  - Categories were derived inductively from reviewing a random sample of 60 reviews before writing any code—ensuring the taxonomy reflected actual user language rather than assumed categories.
- **Step 3.** Prioritization:
  - Scored identified problem clusters on frequency (how many reviews mention it), severity (star rating of reviews mentioning it), and strategic relevance (impact on core user journeys: listing, buying, job-searching).

Dataset Snapshot

Rating	Review Count	% of Total	Response Rate
1-Star (Critical)	546	**61.8%**	71.4%
2-Star (Poor)	74	8.4%	79.7%
3-Star (Neutral)	87	9.8%	83.9%
4-Star (Satisfied)	177	**20.0%**	74.6%

Note: No 5-star reviews were included in this dataset. The analysis is therefore skewed toward active detractors and moderately satisfied users, which is a natural property of review datasets—satisfied users engage with apps, not review stores.

04 | Key Findings

Finding 1: Trust & Safety Is the #1 Pain Point

Across all 884 reviews, trust-related complaints—scams, fraud, and fake content—were the single most recurring theme. 20.0% of all 1-star reviews explicitly mentioned scam, fraud, or fake as their primary grievance. The complaints fell into two distinct sub-patterns:

Trust & Safety — Breakdown of Mentions	
Fake / scam content (general classifieds)	**20.0% of 1-star reviews**
Fake job listings asking for money	**6.4% of 1-star reviews**
Fake property listings (not available)	**Recurring in 2–3-star too**
Account blocked without explanation	**2.6% of 1-star reviews**

Fake job listings are particularly damaging because they attract vulnerable users (job-seekers new to the UAE) who lose both money and trust in the platform. Several reviewers described going to in-person interviews only to be asked for fees of 200–500 AED.

> *"I applied for a job and received so many interview messages but all are asking for money. Very bad. — 1-Star Reviewer"*

> *"Almost all the jobs on this app are fake. You apply and receive so many interviews but all are asking for money. If you are new and looking for a job, avoid this app. — 1-Star Reviewer"*

Finding 2: App Stability & Chat Are Chronic Technical Grievances

Technical reliability complaints were the second most prominent theme, appearing across all star ratings—suggesting these are not isolated incidents but a persistent product quality issue:

Technical Reliability — Breakdown of Mentions	
App crash / not opening / server down	**14.3% of 1-star, 21.6% of 2-star**
Chat / messaging broken or very slow	**8.1% of 1-star, 23.0% of 2-star**
Login / sign-in issues	**2.6% of 1-star reviews**
Filter / search not working	**3.3% overall**
Image / photo upload failing	**Recurring in 2–3-star**

Notably, the chat function was mentioned negatively in 23% of 2-star reviews, making it disproportionately damaging for users who are on the edge of churning. The messaging feature is core to completing transactions on the platform—buyers and sellers need it to close deals—so reliability issues here directly translate to lost revenue for both Dubizzle and its sellers.

> *"I lost many deals because of the defects in chatting. Messages suddenly disappear, get back, get scrambled. Dubizzle is a nice app, but please work on it. — 1-Star Reviewer"*

Finding 3: Response Strategy Is High-Volume but Low-Quality

Dubizzle responds to approximately 74% of reviews, which is notably high compared to industry norms for mobile app support. However, a deeper look at the content of those replies reveals a significant quality gap:

Response Strategy — Quality Analysis	
Total reviews with a Dubizzle reply	**654 out of 884 (74%)**
Replies that redirect to customersupport@dubizzle.com	**78.6% of all replies**
Replies that are exact or near-exact duplicates	**~48% of all replies**
1-star reviews that received NO reply at all	**156 reviews (28.6%)**
Replies addressing the specific complaint raised	**Minority of responses**

The most common reply by volume simply thanks users and directs them to the support email—posted 53 times verbatim in positive (4-star) reviews. Meanwhile, 156 genuinely critical 1-star reviews received no response at all. This pattern signals that the response effort is volume-optimized rather than impact-optimized: the team is replying to easy (positive) reviews while leaving harder (negative, detailed) ones unanswered.

> *“Response example (verbatim, posted 53 times): "Thank you for your review! If you have any suggestions on how we could improve and enhance your user experience, please don't hesitate to share them with us on customersupport@dubizzle.com."”*

Finding 4: Core Value Proposition Still Resonates with Satisfied Users

Among 3 and 4-star reviewers, the platform's fundamental utility remains intact. Satisfied users consistently praise Dubizzle for its coverage and convenience:

- **12.1% of 3–4 star reviews positively mention the buy/sell experience** Buy/sell second-hand goods:
- **11.4% cite it as useful for job search (ironic given 1-star job scam complaints—this signals the problem is execution, not category fit)** Job discovery:
- **52.3% of satisfied reviewers use words like “good,” “great,” “useful,” or “helpful”** Ease of use & general utility:

This is important context: the product is not fundamentally broken. Users who have positive experiences find genuine value. The dissatisfaction is concentrated around specific failure modes, not the platform's core purpose.

05 | Problem Prioritization

Using a simplified Impact × Frequency × Feasibility scoring model, the identified problems were ranked to surface the highest-leverage opportunities. Scores are relative (1–3 scale) and intended to facilitate discussion, not precision.

Problem Area	Frequency	User Impact	Biz Impact	Priority
Fake listings & job scams	★★★	★★★	★★★	P1 — Critical
Chat / messaging reliability	★★★	★★★	★★★	P1 — Critical
App crash / stability	★★★	★★★	★★	P1 — Critical
Response strategy quality	★★★	★★	★★	P2 — High
Ad rejection without refund	★★	★★★	★★	P2 — High
Pricing/fees transparency	★★	★★	★★	P2 — High
Dark mode	★	★	★	P3 — Low

Note: "Feasibility" was excluded from this scoring as it requires internal engineering context I don't have access to. In a real PM role, this score would be populated with input from engineering and design leads.

06 | Product Recommendations

The following recommendations target the three P1 problem clusters. Each is scoped to be actionable at the feature or process level without requiring a platform rebuild.

Recommendation 1: Strengthen Trust Infrastructure for Jobs & Listings

#	Recommendation	Rationale (from data)	Success Metric
1A	Introduce mandatory employer/business verification before posting jobs, with a visible "Verified Employer" badge visible to applicants	6.4% of 1-star reviews cite fake job scams; this is a disproportionately high-harm problem (financial loss, trust erosion for new users)	Reduction in job-related 1-star reviews; flagged report rate on job ads decreases by ~20% within 2 quarters
1B	Add a prominent "Report this ad" shortcut within chat, not just on the listing page, to lower the friction of reporting during live interactions	Scam encounters happen during chats, not on listing pages; current report flow requires users to navigate away from the conversation	Increase in ad-report submissions per active user; reduction in scam-related reviews mentioning that "reporting doesn't work"
1C	Run a periodic automated sweep to deactivate listings older than 90 days that show zero engagement (no messages, no views), with a seller re-confirmation prompt	Multiple reviews mention stale listings wasting time; fake or outdated ads reduce browsing efficiency for genuine buyers	Ratio of active-to-stale listings improves; average time-to-first-message on new listings decreases as noise reduces

Recommendation 2: Stabilize the Chat & Messaging Experience

#	Recommendation	Rationale (from data)	Success Metric
2A	Instrument the chat flow with error telemetry: log message delivery failures, load times > 3s, and session drops. Set an internal SLA of < 2s message delivery in 95% of sessions	8.1% of 1-star and 23.0% of 2-star reviews cite chat as their primary grievance—the second-largest technical pain point after crashes	P95 chat load time < 2s; message delivery failure rate < 1%; chat-related 1-star reviews decline in subsequent quarter's sample
2B	Add a lightweight offline/retry indicator when messages fail to send, so users know their message is queued rather than lost	Multiple reviewers describe sending messages they weren't sure were delivered, then losing deals—this is a perception problem, not just a technical one	User-reported "message lost" complaints decrease; session length in chat increases (users stay in the conversation)

Recommendation 3: Shift Response Strategy from Volume to Value

#	Recommendation	Rationale (from data)	Success Metric
3A	Triage inbound reviews into three buckets: (1) Bug/technical, (2) Policy/payment, (3) General feedback. Route critical 1-star reviews (buckets 1 & 2) to a specialized response queue with a 48-hour SLA and personalized replies	28.6% of 1-star reviews received no reply; 48% of all replies are templates. The current strategy prioritizes volume over quality, missing the highest-impact conversations	1-star response rate increases to 85%+; template reply rate decreases to < 25%; public sentiment score (tracked via 3rd-party tool) stabilizes over 2 quarters
3B	Retire or redesign the three most-used template replies to include at least one specific acknowledgment of the issue raised (e.g., reference the chat problem, the fake listing	53 identical replies sent to 4-star reviewers does not move the needle; resources are better allocated toward converting 1-star detractors	Sentiment score of replies (measured via manual sampling of 50 per month) improves; review-update rate (users who edit their review after a reply) increases

	concern, or the verification bug)		

07 | Success Metrics & Measurement Plan

If these recommendations were implemented, the following metrics would signal progress. Targets are set conservatively for a 2-quarter (6-month) horizon:

Metric	Current Baseline	6-Month Target	Why This Metric
1-star review response rate	71.4%	**85%+**	Measures response coverage improvement
Job ad flag/report rate	No baseline available	**↓ 20% QoQ**	Proxy for fake job listing reduction
Trust/scam mentions in 1-star reviews	20.0% of 1-stars	**↓ to ~15–17%**	Tracks trust issue resolution
Chat-related 1-star mentions	8.1% of 1-stars	**↓ to ~5–6%**	Direct product quality signal
Repeated/template reply rate	~48% of replies	**↓ to < 25%**	Measures response quality shift

Important caveat: app store review data is a lagging, self-selected signal. These metrics should be supplemented with internal retention data, in-app NPS surveys, and support ticket volume to get a complete picture of user health.

08 | Trade-offs & Limitations

What this analysis does not capture

- User segments: Reviews do not distinguish between buyers, sellers, job-seekers, or property seekers—each of whom has a different experience and different needs.
- Frequency of use: A power seller frustrated with a 900 AED ad fee is weighted the same as a one-time user who left a short review.
- Business constraints: Any recommendation around trust verification must be weighed against seller acquisition costs—making verification too strict could reduce listing volume, which is Dubizzle's supply-side moat.
- Timing effects: Several app crash reviews cluster around September 2023, suggesting a specific bad release rather than a chronic issue. A PM would need to validate whether crash rates are still elevated or resolved.
- Silent majority: The 70.1% negative ratio reflects who leaves reviews, not necessarily the full user base. Active daily users rarely leave reviews unless something breaks or delights them.

What I would do differently with access to internal data

- Map reviews to user cohorts (new vs. returning, verified vs. unverified) to understand whether new users churn faster after their first negative encounter.
- Correlate scam reports with listing categories and seller verification status to determine whether the Verified badge actually reduces scam risk.

- A/B test the impact of personalized vs. templated reply responses on review update rates.

09 | Key Learnings & Reflections

This exercise reinforced several PM principles that I want to carry into my early career:

- **1.** Data tells you what, not why. The review data clearly showed that 20% of 1-star reviews mention scams, but it does not explain whether Dubizzle lacks detection tools, lacks enforcement, or simply has not prioritized this. Diagnosis requires mixing quantitative signal with qualitative depth interviews.
- **2.** Prioritization is about trade-offs, not rankings. Recommending job ad verification sounds simple on paper, but it directly competes with the platform's goal of maximizing listing supply. Good prioritization acknowledges this tension rather than ignoring it.
- **3.** Response volume is not the same as response value. Dubizzle's 74% reply rate initially looks strong, but the finding that 48% of replies are near-identical templates reveals a strategy optimized for a metric (reply rate) rather than an outcome (user satisfaction or issue resolution).
- **4.** Start with the user, then work backward. Every recommendation in this document is anchored to a specific user complaint rather than a technology or business need. That discipline—staying user-centric even when it's tempting to propose technically elegant solutions—is what I consider the most important habit for a junior PM to build.

This case study was developed as a portfolio project to demonstrate structured product thinking. All data sourced from public Google Play Store reviews.